

# Automaton Auditor — Final Report (Week 2)

## 1) Executive Summary (What I built, why it matters, and what the system proved)

This project is an **AI-driven auditor** that evaluates a target repository and an architectural PDF report using a “Digital Courtroom” model. Instead of producing a single opinion, the system runs multiple specialized “detectives” to extract **grounded evidence**, then sends that evidence to a panel of “judges” (Prosecutor, Defense, Tech Lead) who debate the quality of the submission under a rubric. Finally, a deterministic “Chief Justice” node synthesizes the judges’ opinions into a structured final audit report.

The core value of the system is that it **reduces hallucination risk** by forcing the pipeline to operate on **explicit evidence objects**. Evidence is collected from two places: the code repository (AST/static checks and repository inspection) and the PDF report (semantic chunking and keyword-based forensic search). Judges are instructed to grade **only based on the evidence provided**, and the Chief Justice then computes the final criterion scores in a transparent way.

In testing the system on my own repository (`--mode self`), the pipeline successfully produced:

- deterministic **repo verification** showing LangGraph fan-out/fan-in structure and Typed State reducers,
- PDF chunk extraction and concept matching,
- judge opinions aligned to rubric criterion IDs,
- a final **AuditReport** saved under `audit/report_onself_generated`.

This final report explains the complete architecture, how evidence is generated, how scoring happens, and what I improved through the iterative MinMax feedback loop.

## 2) Architecture Deep Dive + Diagrams (How the whole system works end-to-end)

### 2.1 System model: Digital Courtroom pipeline

The system is implemented as a **LangGraph StateGraph** built on a single shared **AgentState**. The graph has three layers:

## 1. Detective layer (parallel fan-out)

- Repo Detective: clones the repo safely into a sandbox, runs AST/static checks, flags unsafe execution patterns, verifies graph/state architecture.
- Doc Detective: ingests the PDF using a converter, exports to markdown, creates semantic chunks, and searches for key architecture concepts.
- Vision Inspector: scans repository assets for architecture diagrams/images and produces evidence of existing visuals.

## 2. Judicial bench (parallel fan-out)

- Prosecutor: strict and skeptical
- Defense: generous and credit-giving
- Tech Lead: engineering-focused and practical
  - Each judge outputs a **structured** `JudicialOpinion` object per criterion.

## 3. Chief Justice synthesis

- Groups judge opinions by criterion.
- Computes final criterion scores using deterministic logic.
- Generates an `AuditReport` and writes a markdown report artifact.

This structure matters because it creates **separation of concerns**:

- Detectives do *fact finding* (evidence).
- Judges do *interpretation* (opinions).
- Chief Justice does *aggregation* (final decision).

## 2.2 State design: Typed state + reducers (why it's reliable)

The project uses a TypedDict-based `AgentState` so every node reads/writes predictable fields:

- `repo_url, pdf_path`: inputs
- `evidences`: dictionary of evidence lists keyed by node name

- `opinions`: list of judge outputs
- `final_report`: the output object

Crucially, I used **reducers** so parallel nodes don't overwrite each other:

- Evidence reducer: `operator.ior` merges dictionaries  
`({"repo_detective": [...]}) + {"doc_detective": [...]})`
- Opinions reducer: `operator.add` concatenates lists across judges

This ensures the system works correctly under parallel fan-out (multiple nodes running), because the outputs safely merge instead of competing.

## 2.3 Evidence generation: what counts as “proof” in this system?

Evidence is the foundation of the audit. Every evidence item is a strict Pydantic model:

- goal: what we're trying to prove
- found: True/False
- content: snippet/summary
- location: file path or chunk reference
- rationale: why this evidence supports the goal
- confidence: 0.0–1.0

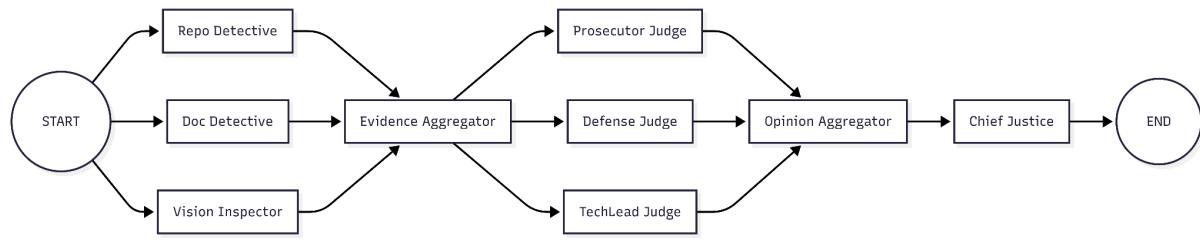
Evidence is produced by:

- static AST checks (`repo_tools`)
- repository scans (security patterns, diagram discovery)
- PDF semantic chunking and targeted search (`doc_tools`)

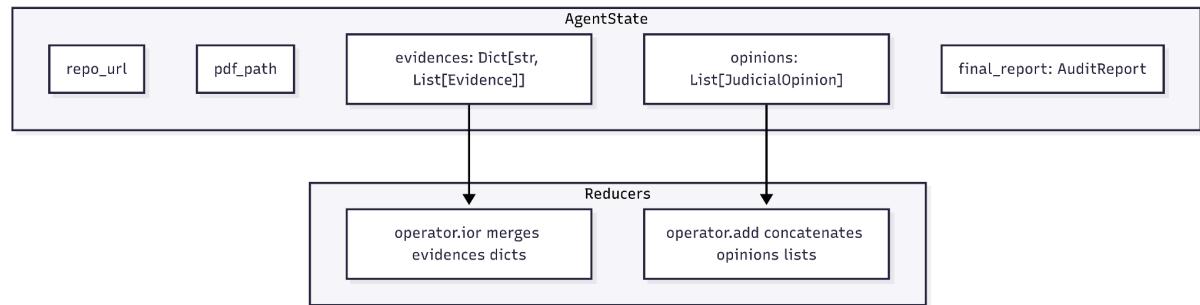
This design makes the whole system auditable: you can trace every claim back to a location.

## 2.4 Diagrams (Mermaid syntax)

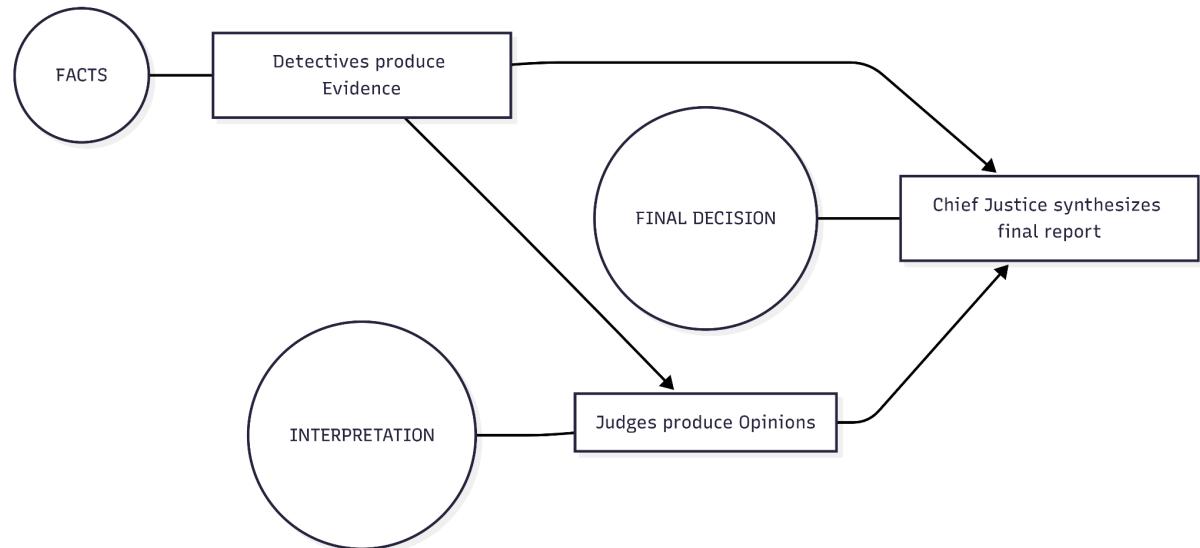
**Diagram A — Full LangGraph pipeline (Detectives → Judges → Chief Justice)**



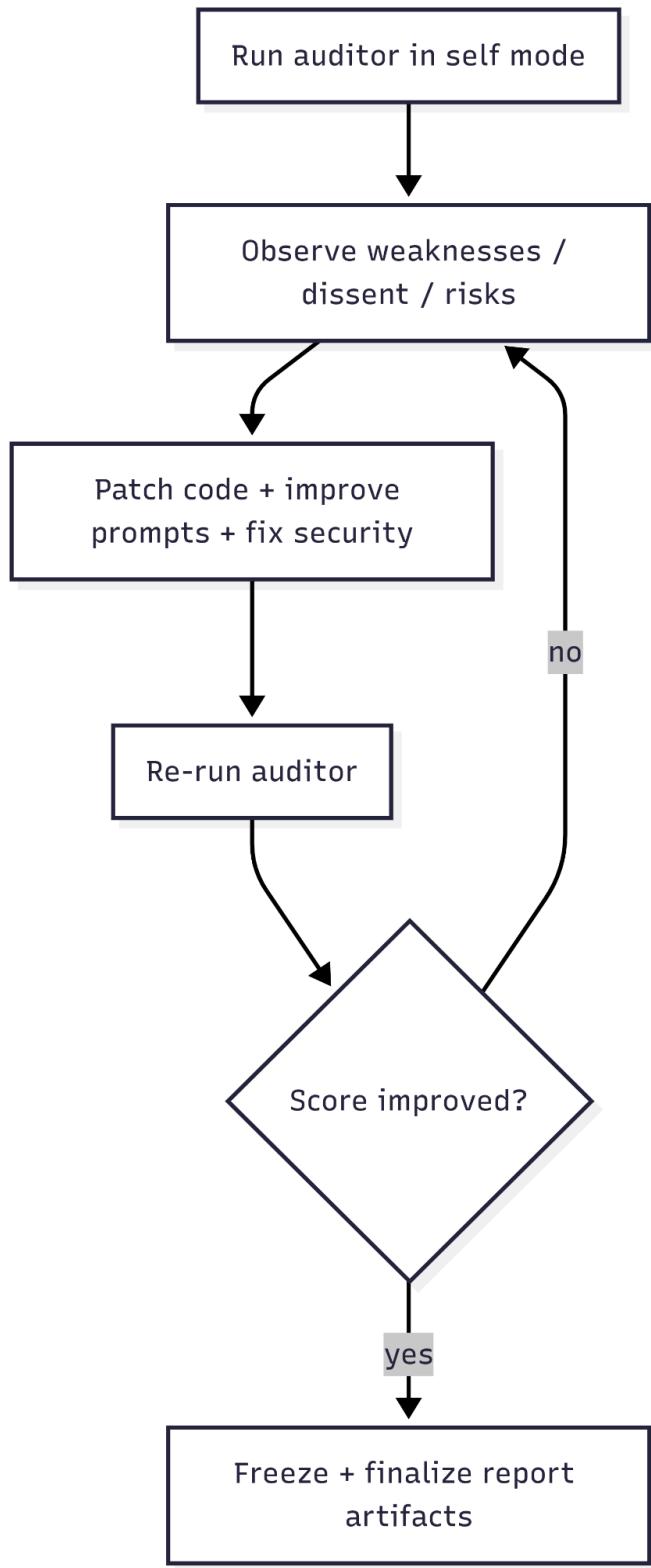
**Diagram B — State + reducers (how parallel nodes merge outputs safely)**



**Diagram C — “Evidence → Opinion → Synthesis” responsibility split**



**Diagram D — MinMax feedback loop (iteration cycle)**



### 3) Self-Audit Criterion Breakdown (How I measured myself and what the results mean)

This section explains how the self-audit is computed and what “grading input” actually is.

### 3.1 What is the “input” to scoring?

The scoring is based on:

1. **Rubric dimensions** loaded from `rubric.json`
2. **Evidence objects** produced by detectives
3. **Judicial opinions** (score + argument + cited\_evidence) produced per criterion
4. **Chief Justice deterministic synthesis logic** aggregating judge scores

So the judges are not grading randomly. They are grading *based on evidence* you collected.

A single criterion is graded like this:

- Evidence exists (or missing)
- Judges see evidence summary + criterion instruction
- Judges output structured score 1..5
- Chief Justice averages scores (with variance logic + dissent)
- Overall score computed from criterion scores

### 3.2 How the rubric is used (and why you shouldn’t “change it”)

Your `rubric.json` is the official rubric from the challenge doc (it uses `dimensions`, `rubric_metadata`, and `synthesis_rules`). That file defines:

- IDs for each dimension (e.g., `forensic_accuracy_code`, `langgraph_architecture`)
- instructions for what to verify
- judge role logic hints
- synthesis rules like security override and dissent requirement

You **should not rewrite this rubric** unless the challenge told you to. Your job is to ensure:

- your detectors collect evidence that maps cleanly to those dimension IDs
- your judges evaluate based on the evidence
- your Chief Justice processes those opinions deterministically

### **3.3 Breakdown of each rubric dimension (what my system checks)**

Below is how your system is aligned to the rubric dimensions you printed (6 dimensions):

#### **(1) `git_forensic_analysis`**

- Goal: assess development progression and commit history quality.
- Evidence source: `extract_git_history(repo_path)` (`repo_tools`)
- What should be improved: add explicit evidence in `repo_investigator` that summarizes commit progression (e.g., number of commits, meaningful messages, chronological evolution). Right now your score is around 3 because the evidence isn't strongly presented to judges.

#### **(2) `security_sandboxing`**

- Goal: verify safe cloning & avoid unsafe execution.
- Evidence source: `clone_repo_sandboxed` (subprocess list, no shell=True), unsafe scan.
- Current status: you fixed unsafe patterns; `unsafe_files: []` is good.
- Improvement: add mention of restricted directories + safe subprocess patterns as explicit evidence.

#### **(3) `forensic_accuracy_code`**

- Goal: verify architecture exists in code (StateGraph, Typed State reducers).
- Evidence source: `verify_graph_forensics()` which checks:
  - `parallel=True` (fan-out)
  - `typed_state=True` (TypedDict + operator.ior)

- Current status: strong evidence exists (score ~4)

#### (4) `forensic_accuracy_docs`

- Goal: verify the PDF's claims match repo reality; detect hallucinations.
- Evidence source: PDF semantic chunks + targeted search for "Dialectical Synthesis", "Metacognition", etc.
- Weakness you saw: score ~3 because your current interim PDF didn't have enough deep detail or explicit cross-references.

#### (5) `judicial_nuance`

- Goal: judges must have distinct prompts + structured output.
- What happened in your run: Prosecutor sometimes scored 1 claiming "no proof of distinct prompts". That means your evidence did not include **proof that judges prompts differ**, and/or there is not a specific evidence item pointing at `src/nodes/judges.py` prompt text.

Fix: In `repo_detective` evidence, add a new evidence item:

- location: `src/nodes/judges.py`
- content: short excerpts showing persona strings differ for Prosecutor/Defense/TechLead
- found=True  
Then Prosecutor cannot claim "no proof".

#### (6) `langgraph_architecture`

- Goal: verify fan-out/fan-in across detectives and judges; reducers; error handling edges.
- You already have the fan-out/fan-in. To reach 5:
  - add conditional edges or explicit failure handling (optional)
  - include evidence demonstrating fan-in aggregator nodes exist
  - include reducers evidence

## 4) MinMax Feedback Loop Reflection (What I changed after feedback and why)

This project naturally forced a “MinMax” iteration: I ran the system, observed its strictest critic (Prosecutor), then improved the weakest areas.

Key iteration improvements:

### 1. Security cleanup

- Early versions were flagged because “unsafe system execution” patterns appeared in evidence.
- I refactored repo cloning to use safe subprocess calls (no `shell=True`).
- I improved the unsafe scanner to skip `.venv`, caches, and self-files to avoid false positives.  
Result: `unsafe_files=[ ]` and security risk reduced.

### 2. Evidence quality improvements

- I changed evidence to be shorter but more explicit (clipped snippets, clear locations).
- I deduplicated PDF findings to prevent multiple chunks causing “inconsistent evidence” complaints.

### 3. Structured output enforcement

- Judges now use `with_structured_output(JudicialOpinion)` so outputs conform to schema.
- This removes messy free-text judge outputs and improves reproducibility.

### 4. Chief Justice fairness

- Instead of hard-capping overall score to 2 when security is flagged, I applied a penalty approach.
- This avoids a “destroy everything” outcome and keeps scoring fair.

The loop taught me something important: **even good code can score low if evidence is not presented clearly**. The auditor is only as strong as its evidence traceability.

---

## 5) Remediation Plan (Exactly what I'll improve next, in order)

This remediation plan is focused on raising the system's self-audit scores and also making the auditor useful for peer grading.

### Priority 1 — Fix judicial\_nuance score (make judges prompts provable)

Add explicit repo evidence:

- Read `src/nodes/judges.py`
- Extract 2–3 short lines from each persona prompt
- Create evidence objects proving prompt differences  
This directly blocks the Prosecutor's "persona collusion" accusation.

### Priority 2 — Improve git\_forensic\_analysis evidence

Add a dedicated evidence object summarizing:

- total commits
- first and last commit messages
- progression pattern (setup → detection → judging → synthesis)  
That will push the score from ~3 to 4+.

### Priority 3 — Strengthen forensic\_accuracy\_docs with cross-references

In the PDF, add:

- a section that explicitly says "Claim: parallel judges implemented in graph.py lines X–Y"
- include screenshots or code snippets and references  
This helps the Prosecutor validate "no hallucination".

### Priority 4 — Extend langgraph\_architecture robustness

Optional improvements for maximum score:

- Add conditional edges or failure handling (e.g., if repo cloning fails, bypass some nodes).
- Add explicit state flags like `repo_ok`, `pdf_ok`.  
This makes the orchestration “rigor” stronger.

## Priority 5 — Improve output artifacts

Ensure every run creates:

- JSON final report
- Markdown audit report
- saved folder under `audit/report_*`
- include LangSmith trace screenshot/link if required by submission